

# Lecture 6: Recursive Nonlinear Non-Gaussian Estimation, Part II

## AER1513: State Estimation

Timothy D. Barfoot

University of Toronto

Copyright © 2023

# Outline

## Lecture 6: Recursive Nonlinear Non-Gaussian Estimation, Part II

- Motivation and Recap

- Problem Setup

- Filter Taxonomy

- Transformation Methods

- Particle Filter

- Sigmapoint Kalman Filter

# Motivation

- last lecture we looked at recursive estimation techniques for nonlinear non-Gaussian situations
- we derived the Bayes filter, generalized Gaussian filter, and finally the IEKF and EKF
- this lecture, we will continue to explore some other options in our filter taxonomy that avoid the use of linearization
- we'll introduce the **sigmapoint transformation** as a replacement to linearization, and use it to derive the **sigmapoint Kalman filter**
- we'll then return to the Bayes filter and branch off in a different way to derive the **particle filter**

# System

- we define our system using the following **nonlinear, time-varying** models:

**motion model:**  $\mathbf{x}_k = \mathbf{f}(\mathbf{x}_{k-1}, \mathbf{v}_k, \mathbf{w}_k), \quad k = 1 \dots K \quad (1a)$

**observation model:**  $\mathbf{y}_k = \mathbf{g}(\mathbf{x}_k, \mathbf{n}_k), \quad k = 0 \dots K \quad (1b)$

where  $k$  is again the discrete-time index and  $K$  its maximum

- the variables have the following meanings:

system state :  $\mathbf{x}_k \in \mathbb{R}^N$

initial state :  $\mathbf{x}_0 \in \mathbb{R}^N \sim \mathcal{N}(\check{\mathbf{x}}_0, \check{\mathbf{P}}_0)$

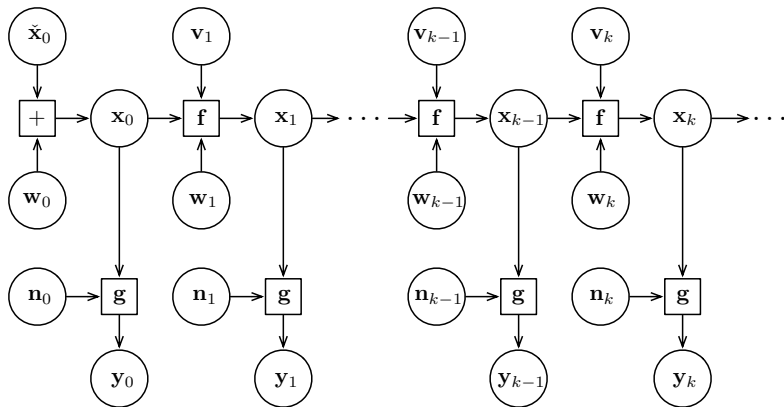
input :  $\mathbf{v}_k \in \mathbb{R}^N$

process noise :  $\mathbf{w}_k \in \mathbb{R}^N \sim \mathcal{N}(\mathbf{0}, \mathbf{Q}_k)$

measurement :  $\mathbf{y}_k \in \mathbb{R}^M$

measurement noise :  $\mathbf{n}_k \in \mathbb{R}^M \sim \mathcal{N}(\mathbf{0}, \mathbf{R}_k)$

# Markov property



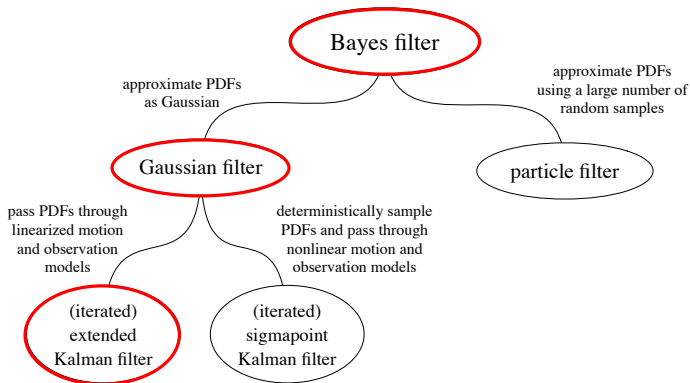
# Problem statement

- our **state estimation problem** is still the following:

## Definition

The problem of **state estimation** is to come up with an estimate,  $\hat{\mathbf{x}}_k$ , of the true *state* of a system, at one or more timesteps,  $k$ , given knowledge of the initial state,  $\check{\mathbf{x}}_0$ , a sequence of measurements,  $\mathbf{y}_{0:K,\text{meas}}$ , a sequence of inputs,  $\mathbf{v}_{1:K}$ , as well as knowledge of the system's motion and observation models.

# Filter taxonomy



# Bayes filter

- recall the very general **Bayes filter**:

$$\underbrace{p(\mathbf{x}_k | \check{\mathbf{x}}_0, \mathbf{v}_{1:k}, \mathbf{y}_{0:k})}_{\text{posterior belief}}$$
$$= \eta \underbrace{p(\mathbf{y}_k | \mathbf{x}_k)}_{\substack{\text{observation} \\ \text{correction} \\ \text{using } \mathbf{g}(\cdot)}} \int \underbrace{p(\mathbf{x}_k | \mathbf{x}_{k-1}, \mathbf{v}_k)}_{\substack{\text{motion} \\ \text{prediction} \\ \text{using } \mathbf{f}(\cdot)}} \underbrace{p(\mathbf{x}_{k-1} | \check{\mathbf{x}}_0, \mathbf{v}_{1:k-1}, \mathbf{y}_{0:k-1})}_{\text{prior belief}} d\mathbf{x}_{k-1}$$

(2)

- we can see this takes on a **predictor-corrector** form, just like the KF



# Generalized Gaussian filter

- assuming all the densities are Gaussian we have the **generalized Gaussian filter**:

**predictor:**  $\check{\mathbf{P}}_k = \mathbf{F}_{k-1} \hat{\mathbf{P}}_{k-1} \mathbf{F}_{k-1}^T + \mathbf{Q}'_k$  (3a)

$$\check{\mathbf{x}}_k = \mathbf{f}(\hat{\mathbf{x}}_{k-1}, \mathbf{v}_k, \mathbf{0})$$
 (3b)

**Kalman gain:**  $\mathbf{K}_k = \Sigma_{xy,k} \Sigma_{yy,k}^{-1}$  (3c)

$$\hat{\mathbf{P}}_k = \check{\mathbf{P}}_k - \mathbf{K}_k \Sigma_{xy,k}^T$$
 (3d)

**corrector:**  $\hat{\mathbf{x}}_k = \check{\mathbf{x}}_k + \mathbf{K}_k (\mathbf{y}_k - \boldsymbol{\mu}_{y,k})$  (3e)

- in this case, the moments come from constructing the joint Gaussian (state and measurement)...somehow

$$p(\mathbf{x}_k, \mathbf{y}_k | \check{\mathbf{x}}_0, \mathbf{v}_{1:k}, \mathbf{y}_{0:k-1}) = \mathcal{N} \left( \begin{bmatrix} \boldsymbol{\mu}_{x,k} \\ \boldsymbol{\mu}_{y,k} \end{bmatrix}, \begin{bmatrix} \Sigma_{xx,k} & \Sigma_{xy,k} \\ \Sigma_{yx,k} & \Sigma_{yy,k} \end{bmatrix} \right) \quad (4)$$

# IEKF

- using linearization to compute the statistical moment in the joint Gaussian we have

**predictor:**

$$\check{\mathbf{P}}_k = \mathbf{F}_{k-1} \hat{\mathbf{P}}_{k-1} \mathbf{F}_{k-1}^T + \mathbf{Q}'_k \quad (5a)$$

$$\check{\mathbf{x}}_k = \mathbf{f}(\hat{\mathbf{x}}_{k-1}, \mathbf{v}_k, \mathbf{0}) \quad (5b)$$

**Kalman gain:**

$$\mathbf{K}_k = \check{\mathbf{P}}_k \mathbf{G}_k^T (\mathbf{G}_k \check{\mathbf{P}}_k \mathbf{G}_k^T + \mathbf{R}'_k)^{-1} \quad (5c)$$

$$\hat{\mathbf{P}}_k = (\mathbf{1} - \mathbf{K}_k \mathbf{G}_k) \check{\mathbf{P}}_k \quad (5d)$$

**corrector:**

$$\hat{\mathbf{x}}_k = \check{\mathbf{x}}_k + \mathbf{K}_k (\mathbf{y}_k - \mathbf{y}_{\text{op},k} - \mathbf{G}_k(\check{\mathbf{x}}_k - \mathbf{x}_{\text{op},k})) \quad (5e)$$

- in this case, we iterate the last three equations to convergence
- the ‘mean’ of the IEKF actually converges to the mode of the posterior

# EKF

- finally, with just one iteration we have the EKF

predictor:  $\check{\mathbf{P}}_k = \mathbf{F}_{k-1} \hat{\mathbf{P}}_{k-1} \mathbf{F}_{k-1}^T + \mathbf{Q}'_k$  (6a)


$$\check{\mathbf{x}}_k = \mathbf{f}(\hat{\mathbf{x}}_{k-1}, \mathbf{v}_k, \mathbf{0})$$
 (6b)

Kalman gain:  $\mathbf{K}_k = \check{\mathbf{P}}_k \mathbf{G}_k^T (\mathbf{G}_k \check{\mathbf{P}}_k \mathbf{G}_k^T + \mathbf{R}'_k)^{-1}$  (6c)

$$\hat{\mathbf{P}}_k = (\mathbf{1} - \mathbf{K}_k \mathbf{G}_k) \check{\mathbf{P}}_k$$
 (6d)


corrector:  $\hat{\mathbf{x}}_k = \check{\mathbf{x}}_k + \mathbf{K}_k (\mathbf{y}_k - \mathbf{g}(\check{\mathbf{x}}_k, \mathbf{0}))$  (6e)

# Transforming PDFs through nonlinearities




- we've seen how to use **linearization** to transform a Gaussian through a nonlinearity
- there are some other techniques we can use:
  - **Monte Carlo sampling** can be used to transform any PDF through a nonlinearity
  - **sigmapoint transformation** can be used to transform a Gaussian through a nonlinearity

# Monte Carlo sampling



- **Monte Carlo** just draws many random samples from the input, transforms each sample, then recombines the samples (e.g., compute moments)
- it's very inefficient, works with any type of PDF, any type of nonlinearity

# Sigmapoint transformation



- the **sigmapoint transformation** draws a small number of deterministic samples, transforms each sample, then recombines using special moment formulas
- it's quite efficient, works with Gaussian PDFs

## SP: step 1

- a set of  $2L + 1$  **sigmapoints** is computed from the input density,  $\mathcal{N}(\boldsymbol{\mu}_x, \boldsymbol{\Sigma}_{xx})$ , according to

$$\mathbf{L}\mathbf{L}^T = \boldsymbol{\Sigma}_x \quad (\text{Cholesky decomposition, } \mathbf{L} \text{ lower-triangular}) \quad (7a)$$


$$\mathbf{x}_0 = \boldsymbol{\mu}_x \quad (7b)$$

$$\mathbf{x}_i = \boldsymbol{\mu}_x + \sqrt{L + \kappa} \text{col}_i \mathbf{L} \quad (7c)$$

$$\mathbf{x}_{i+L} = \boldsymbol{\mu}_x - \sqrt{L + \kappa} \text{col}_i \mathbf{L} \quad i = 1 \dots L \quad (7d)$$

where  $L = \dim(\boldsymbol{\mu}_x)$

# SP: step 1





## SP: step 1

- the original moments can easily be rebuilt from the sigmapoints:

$$\boldsymbol{\mu}_x = \sum_{i=0}^{2L} \alpha_i \mathbf{x}_i \quad (8a)$$

$$\boldsymbol{\Sigma}_{xx} = \sum_{i=0}^{2L} \alpha_i (\mathbf{x}_i - \boldsymbol{\mu}_x) (\mathbf{x}_i - \boldsymbol{\mu}_x)^T \quad (8b)$$

where

$$\alpha_i = \begin{cases} \frac{\kappa}{L+\kappa} & i = 0 \\ \frac{1}{2} \frac{1}{L+\kappa} & \text{otherwise} \end{cases} \quad (9)$$

which we note sums to 1

- the user-definable parameter,  $\kappa$ , will be explained later

## SP: steps 2-5

- each of the sigmapoints is individually passed through the nonlinearity,  $\mathbf{g}(\cdot)$ :

$$\mathbf{y}_i = \mathbf{g}(\mathbf{x}_i), \quad i = 0 \dots 2L \quad (10)$$

- the **mean** of the output density,  $\boldsymbol{\mu}_y$ , is computed as

$$\boldsymbol{\mu}_y = \sum_{i=0}^{2L} \alpha_i \mathbf{y}_i \quad (11)$$

- the **covariance** of the output density,  $\boldsymbol{\Sigma}_y$ , is computed as

$$\boldsymbol{\Sigma}_{yy} = \sum_{i=0}^{2L} \alpha_i (\mathbf{y}_i - \boldsymbol{\mu}_y) (\mathbf{y}_i - \boldsymbol{\mu}_y)^T \quad (12)$$

- the output density,  $\mathcal{N}(\boldsymbol{\mu}_y, \boldsymbol{\Sigma}_{yy})$ , is returned

# SP commentary

- by approximating the input density instead of linearizing, we avoid the need to compute the Jacobian of the nonlinearity (either in closed form or numerically)
- we employ only standard linear algebra operations (Cholesky decomposition, outer products, matrix summations)
- the computation cost is similar to linearization (when a numerical Jacobian is used)
- there is no requirement that the nonlinearity be smooth and differentiable

# Comparison example

- we'll use a simple example to compare the various transformation methods:
  - Monte Carlo
  - linearization
  - sigmapoint
- consider the simple one-dimensional nonlinearity

$$f(x) = x^2 \tag{13}$$

- let the prior density be

$$\mathcal{N}(\mu_x, \sigma_x^2) \tag{14}$$

## Example: Monte Carlo

- for this example, we can work out the Monte Carlo method analytically
- an arbitrary sample (a.k.a., realization) of the input density is given by

$$x_i = \mu_x + \delta x_i, \quad \delta x_i \leftarrow \mathcal{N}(0, \sigma_x^2) \quad (15)$$

- transforming this sample through the nonlinearity we get

$$y_i = f(x_i) = f(\mu_x + \delta x_i) = (\mu_x + \delta x_i)^2 = \mu_x^2 + 2\mu_x \delta x_i + \delta x_i^2 \quad (16)$$

- taking the expectation of both sides we arrive at the **mean** of the output:

$$\mu_y = E[y_i] = \mu_x^2 + 2\mu_x \underbrace{E[\delta x_i]}_0 + \underbrace{E[\delta x_i^2]}_{\sigma_x^2} = \mu_x^2 + \sigma_x^2 \quad (17)$$

## Example: Monte Carlo

- we do a similar thing for the **variance** of the output:

$$\sigma_y^2 = E[(y_i - \mu_y)^2] \quad (18a)$$

$$= E[(2\mu_x \delta x_i + \delta x_i^2 - \sigma_x^2)^2] \quad (18b)$$

$$= \underbrace{E[\delta x_i^4]}_{3\sigma_x^4} + 4\mu_x \underbrace{E[\delta x_i^3]}_0 + (4\mu_x^2 - 2\sigma_x^2) \underbrace{E[\delta x_i^2]}_{\sigma_x^2} - 4\mu_x \sigma_x^2 \underbrace{E[\delta x_i]}_0 + \sigma_x^4 \quad (18c)$$

$$= 4\mu_x^2 \sigma_x^2 + 2\sigma_x^4 \quad (18d)$$

where  $E[\delta x_i^3] = 0$  and  $E[\delta x_i^4] = 3\sigma_x^4$  are the well-known third and fourth moments for a Gaussian PDF

## Example: linearization

- linearizing about the mean of the input density we have

$$y_i = f(\mu_x + \delta x_i) \approx \underbrace{f(\mu_x)}_{\mu_x^2} + \underbrace{\frac{\partial f}{\partial x} \Big|_{\mu_x}}_{2\mu_x} \delta x_i = \mu_x^2 + 2\mu_x \delta x_i \quad (19)$$

- taking the expectation we arrive at the mean of the output:

$$\mu_y = E[y_i] = \mu_x^2 + 2\mu_x \underbrace{E[\delta x_i]}_0 = \mu_x^2 \quad (20)$$

- for the variance of the output we have

$$\sigma_y^2 = E[(y_i - \mu_y)^2] = E[(2\mu_x \delta x_i)^2] = 4\mu_x^2 \sigma_x^2 \quad (21)$$

- this estimate is **biased** and **overconfident** compared to Monte Carlo

## Example: sigmapoint

- there are  $2L + 1 = 3$  **sigmapoints** in dimension  $L = 1$ :

$$x_0 = \mu_x, \quad x_1 = \mu_x + \sqrt{1 + \kappa} \sigma_x, \quad x_2 = \mu_x - \sqrt{1 + \kappa} \sigma_x \quad (22)$$

where  $\kappa$  is a user-definable parameter that we discuss below

- we pass each sigmapoint through the nonlinearity:

$$y_0 = f(x_0) = \mu_x^2 \quad (23a)$$

$$\begin{aligned} y_1 &= f(x_1) = (\mu_x + \sqrt{1 + \kappa} \sigma_x)^2 \\ &= \mu_x^2 + 2\mu_x \sqrt{1 + \kappa} \sigma_x + (1 + \kappa) \sigma_x^2 \end{aligned} \quad (23b)$$

$$\begin{aligned} y_2 &= f(x_2) = (\mu_x - \sqrt{1 + \kappa} \sigma_x)^2 \\ &= \mu_x^2 - 2\mu_x \sqrt{1 + \kappa} \sigma_x + (1 + \kappa) \sigma_x^2 \end{aligned} \quad (23c)$$



## Example: sigmapoint

- the **mean** of the output is given by

$$\mu_y = \frac{1}{1 + \kappa} \left( \kappa y_0 + \frac{1}{2} \sum_{i=1}^2 y_i \right) \quad (24a)$$

$$= \frac{1}{1 + \kappa} \left( \kappa \mu_x^2 + \frac{1}{2} (\mu_x^2 + 2\mu_x \sqrt{1 + \kappa} \sigma_x + (1 + \kappa) \sigma_x^2 + \mu_x^2 - 2\mu_x \sqrt{1 + \kappa} \sigma_x + (1 + \kappa) \sigma_x^2) \right) \quad (24b)$$

$$= \frac{1}{1 + \kappa} (\kappa \mu_x^2 + \mu_x^2 + (1 + \kappa) \sigma_x^2) \quad (24c)$$

$$= \mu_x^2 + \sigma_x^2 \quad (24d)$$

which is independent of  $\kappa$  and exactly the same as Monte Carlo

## Example: sigmapoint

- for the **variance** we have

$$\sigma_y^2 = \frac{1}{1 + \kappa} \left( \kappa (y_0 - \mu_y)^2 + \frac{1}{2} \sum_{i=1}^2 (y_i - \mu_y)^2 \right) \quad (25a)$$

$$= \frac{1}{1 + \kappa} \left( \kappa \sigma_x^4 + \frac{1}{2} \left( (2\mu_x \sqrt{1 + \kappa} \sigma_x + \kappa \sigma_x^2)^2 + (-2\mu_x \sqrt{1 + \kappa} \sigma_x + \kappa \sigma_x^2)^2 \right) \right) \quad (25b)$$

$$= \frac{1}{1 + \kappa} (\kappa \sigma_x^4 + 4(1 + \kappa) \mu_x^2 \sigma_x^2 + \kappa^2 \sigma_x^4) \quad (25c)$$

$$= 4\mu_x^2 \sigma_x^2 + \kappa \sigma_x^4 \quad (25d)$$

which can be made to be identical to Monte Carlo by selecting the user-definable parameter,  $\kappa$ , to be 2

- thus, for this nonlinearity, the unscented transformation can exactly capture the correct mean and variance of the output

## Example: sigmapoint

- we choose  $\kappa$  to make the fourth moment of the sigmapoints match the true **kurtosis** of the Gaussian input density:

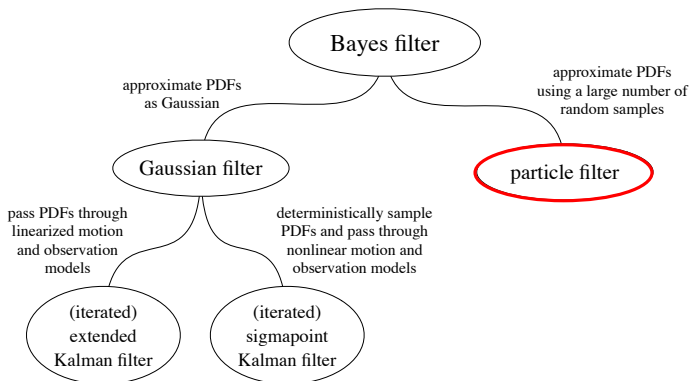
$$3\sigma_x^4 = \frac{1}{1+\kappa} \left( \underbrace{\kappa (x_0 - \mu_x)^4}_0 + \frac{1}{2} \sum_{i=1}^2 (x_i - \mu_x)^4 \right) \quad (26a)$$

$$= \frac{1}{2(1+\kappa)} \left( (\sqrt{1+\kappa}\sigma_x)^4 + (-\sqrt{1+\kappa}\sigma_x)^4 \right) \quad (26b)$$

$$= (1 + \kappa)\sigma_x^4 \quad (26c)$$

- comparing the desired and actual kurtosis, we should pick  $\kappa = 2$  to make them match exactly
- not surprisingly, this has a positive effect on accuracy of the transformation

# Particle filter



# Particle filter

- let's try to use the **Monte Carlo** transformation method to implement the Bayes filter
- the **particle filter** is one of the only practical techniques able to handle non-Gaussian noise and nonlinear observation and motion models
- it is practical in that it is very easy to implement; we do not even need to have analytical expressions for  $f(\cdot)$  and  $g(\cdot)$ , nor their derivatives
- can be inefficient due to the need for many samples
- sometimes called the **bootstrap algorithm**, the **condensation algorithm**, or the **survival-of-the-fittest algorithm**

## PF: prediction

- draw  $M$  **samples** from the joint density comprising the prior and the motion noise:

$$\begin{bmatrix} \hat{\mathbf{x}}_{k-1,m} \\ \mathbf{w}_{k,m} \end{bmatrix} \leftarrow p(\mathbf{x}_{k-1} | \check{\mathbf{x}}_0, \mathbf{v}_{1:k-1}, \mathbf{y}_{1:k-1}) p(\mathbf{w}_k) \quad (27)$$

where  $m$  is the unique particle index

- generate a prediction of the posterior PDF by using  $\mathbf{v}_k$
- this is done by passing each prior particle/noise sample through the nonlinear motion model:

$$\check{\mathbf{x}}_{k,m} = \mathbf{f}(\hat{\mathbf{x}}_{k-1,m}, \mathbf{v}_k, \mathbf{w}_{k,m}) \quad (28)$$

- these new **predicted particles** together approximate the density,  $p(\mathbf{x}_k | \check{\mathbf{x}}_0, \mathbf{v}_{1:k}, \mathbf{y}_{1:k-1})$

## PF: correction

- correct the posterior PDF by incorporating  $\mathbf{y}_k$
- first, assign a scalar **weight**,  $w_{k,m}$ , to each predicted particle based on the **divergence** between the desired posterior and the predicted posterior for each particle:

$$w_{k,m} = \frac{p(\check{\mathbf{x}}_{k,m} | \check{\mathbf{x}}_0, \mathbf{v}_{1:k}, \mathbf{y}_{1:k})}{p(\check{\mathbf{x}}_{k,m} | \check{\mathbf{x}}_0, \mathbf{v}_{1:k}, \mathbf{y}_{1:k-1})} = \eta p(\mathbf{y}_k | \check{\mathbf{x}}_{k,m}) \quad (29)$$

where  $\eta$  is a normalization constant

- this is typically accomplished in practice by simulating an expected sensor reading,  $\check{\mathbf{y}}_{k,m}$ , using the nonlinear observation model:

$$\check{\mathbf{y}}_{k,m} = \mathbf{g}(\check{\mathbf{x}}_{k,m}, \mathbf{0}) \quad (30)$$

- we then assume  $p(\mathbf{y}_k | \check{\mathbf{x}}_{k,m}) = p(\mathbf{y}_k | \check{\mathbf{y}}_{k,m})$ , where the right-hand side is a known density (e.g., Gaussian)

## PF: correction

- second, **resample** the posterior based on the weight assigned to each predicted posterior particle:

$$\hat{\mathbf{x}}_{k,m} \xleftarrow{\text{resample}} \{\check{\mathbf{x}}_{k,m}, w_{k,m}\} \quad (31)$$

- this can be done in several different ways (e.g., Madow systematic sampling)
- create  $M$  **bins** based on weights

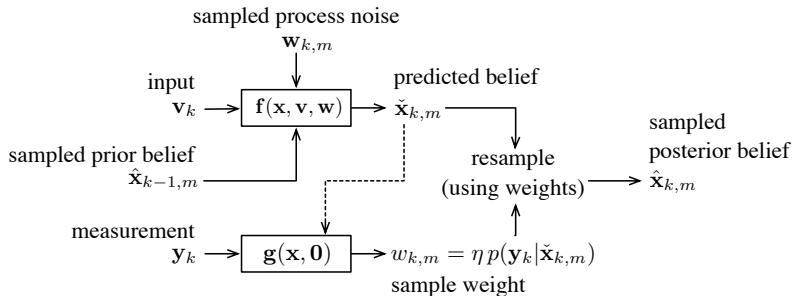
$$\beta_m = \frac{\sum_{n=1}^m w_n}{\sum_{\ell=1}^M w_\ell} \quad (32)$$

$$0 \leq \beta_1 \leq \beta_2 \leq \dots \leq \beta_{M-1} \leq 1$$

- march along in steps of  $1/M$ , adding a sample from each bin that is visited



# PF in pictures

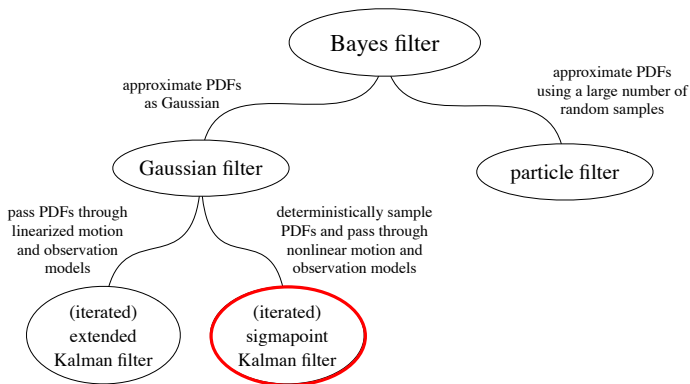


# PF in action


# PF commentary

- we do not necessarily need to resample every time we go through the algorithm
- to be on the safe side, it is wise to add a small percentage of samples in the prediction step that are uniformly drawn from the entire state sample space
- for high-dimensional state estimation problems, the PF can become computationally intractable
- the PF is an **anytime** algorithm – we can just keep adding particles until we run out of time, then resample and give an answer
- we can dynamically pick the number of particles online using a heuristic such as  $\sum w_{k,m} \geq w_{\text{thresh}}$ , a threshold
- the CRLB is set by the uncertainty in the measurements that we have available – using more samples does not allow us to do better

# Sigmapoint Kalman filter



- we can also derive a filtering scheme that uses the **sigmapoint transformation** to implement the **generalized Gaussian filter**, rather than **linearization**



## SPKF: prediction

- both the prior belief and the motion noise have uncertainty so these are stacked together in the following way:

$$\boldsymbol{\mu}_z = \begin{bmatrix} \hat{\mathbf{x}}_{k-1} \\ \mathbf{0} \end{bmatrix}, \quad \boldsymbol{\Sigma}_{zz} = \begin{bmatrix} \hat{\mathbf{P}}_{k-1} & \mathbf{0} \\ \mathbf{0} & \mathbf{Q}_k \end{bmatrix} \quad (33)$$

where we see that  $\{\boldsymbol{\mu}_z, \boldsymbol{\Sigma}_{zz}\}$  is still a Gaussian representation

- we let  $L = \dim \boldsymbol{\mu}_z$
- convert  $\{\boldsymbol{\mu}_z, \boldsymbol{\Sigma}_{zz}\}$  to a **sigmapoint** representation:

$$\mathbf{L}\mathbf{L}^T = \boldsymbol{\Sigma}_{zz} \quad (\text{Cholesky decomposition, } \mathbf{L} \text{ lower-triangular}) \quad (34a)$$

$$\mathbf{z}_0 = \boldsymbol{\mu}_z \quad (34b)$$

$$\mathbf{z}_i = \boldsymbol{\mu}_z + \sqrt{L + \kappa} \operatorname{col}_i \mathbf{L} \quad (34c)$$

$$\mathbf{z}_{i+L} = \boldsymbol{\mu}_z - \sqrt{L + \kappa} \operatorname{col}_i \mathbf{L} \quad i = 1 \dots L \quad (34d)$$

# SPKF: prediction

- unstack each **sigmapoint** into state and motion noise,

$$\mathbf{z}_i = \begin{bmatrix} \hat{\mathbf{x}}_{k-1,i} \\ \mathbf{w}_{k,i} \end{bmatrix} \quad (35)$$

- pass each sigmapoint through the nonlinear motion model exactly:

$$\check{\mathbf{x}}_{k,i} = \mathbf{f}(\hat{\mathbf{x}}_{k-1,i}, \mathbf{v}_k, \mathbf{w}_{k,i}) \quad i = 0 \dots 2L \quad (36)$$

- the latest input,  $\mathbf{v}_k$ , is required

## SPKF: prediction

- recombine the transformed sigmapoints into the **predicted belief**,  $\{\check{\mathbf{x}}_k, \check{\mathbf{P}}_k\}$ , according to

$$\check{\mathbf{x}}_k = \sum_{i=0}^{2L} \alpha_i \check{\mathbf{x}}_{k,i} \quad (37a)$$

$$\check{\mathbf{P}}_k = \sum_{i=0}^{2L} \alpha_i (\check{\mathbf{x}}_{k,i} - \check{\mathbf{x}}_k) (\check{\mathbf{x}}_{k,i} - \check{\mathbf{x}}_k)^T \quad (37b)$$

where

$$\alpha_i = \begin{cases} \frac{\kappa}{L+\kappa} & i = 0 \\ \frac{1}{2} \frac{1}{L+\kappa} & \text{otherwise} \end{cases} \quad (38)$$



# SPKF: correction

- this step is a little more subtle
- we'll begin by recalling the **generalized Gaussian** correction step:

$$\mathbf{K}_k = \Sigma_{xy,k} \Sigma_{yy,k}^{-1} \quad (39a)$$

$$\hat{\mathbf{P}}_k = \check{\mathbf{P}}_k - \mathbf{K}_k \Sigma_{xy,k}^T \quad (39b)$$

$$\hat{\mathbf{x}}_k = \check{\mathbf{x}}_k + \mathbf{K}_k (\mathbf{y}_k - \boldsymbol{\mu}_{y,k}) \quad (39c)$$

- we'll use the **sigmapoint transformation** to come up with better versions of  $\boldsymbol{\mu}_{y,k}$ ,  $\Sigma_{yy,k}$ , and  $\Sigma_{xy,k}$  than we found through linearization

# SPKF: correction

- both the predicted belief and the observation noise have uncertainty so these are stacked together in the following way:

$$\boldsymbol{\mu}_z = \begin{bmatrix} \check{\mathbf{x}}_k \\ \mathbf{0} \end{bmatrix}, \quad \boldsymbol{\Sigma}_{zz} = \begin{bmatrix} \check{\mathbf{P}}_k & \mathbf{0} \\ \mathbf{0} & \mathbf{R}_k \end{bmatrix} \quad (40)$$

where we see that  $\{\boldsymbol{\mu}_z, \boldsymbol{\Sigma}_{zz}\}$  is still a Gaussian representation

- we let  $L = \dim \boldsymbol{\mu}_z$
- convert  $\{\boldsymbol{\mu}_z, \boldsymbol{\Sigma}_{zz}\}$  to a **sigmapoint representation**:

$$\mathbf{L}\mathbf{L}^T = \boldsymbol{\Sigma}_{zz} \quad (\text{Cholesky decomposition, } \mathbf{L} \text{ lower-triangular}) \quad (41a)$$

$$\mathbf{z}_0 = \boldsymbol{\mu}_z \quad (41b)$$

$$\mathbf{z}_i = \boldsymbol{\mu}_z + \sqrt{L + \kappa} \text{col}_i \mathbf{L} \quad (41c)$$

$$\mathbf{z}_{i+L} = \boldsymbol{\mu}_z - \sqrt{L + \kappa} \text{col}_i \mathbf{L} \quad i = 1 \dots L \quad (41d)$$

# SPKF: correction

- unstack each sigmapoint into state and observation noise,

$$\mathbf{z}_i = \begin{bmatrix} \check{\mathbf{x}}_{k,i} \\ \mathbf{n}_{k,i} \end{bmatrix} \quad (42)$$

- pass each sigmapoint through the nonlinear observation model exactly:

$$\check{\mathbf{y}}_{k,i} = \mathbf{g}(\check{\mathbf{x}}_{k,i}, \mathbf{n}_{k,i}) \quad (43)$$

# SPKF: correction

- recombine the transformed sigmapoints into the **desired moments**:

$$\boldsymbol{\mu}_{y,k} = \sum_{i=0}^{2L} \alpha_i \check{\mathbf{y}}_{k,i} \quad (44a)$$

$$\boldsymbol{\Sigma}_{yy,k} = \sum_{i=0}^{2L} \alpha_i (\check{\mathbf{y}}_{k,i} - \boldsymbol{\mu}_{y,k}) (\check{\mathbf{y}}_{k,i} - \boldsymbol{\mu}_{y,k})^T \quad (44b)$$

$$\boldsymbol{\Sigma}_{xy,k} = \sum_{i=0}^{2L} \alpha_i (\check{\mathbf{x}}_{k,i} - \check{\mathbf{x}}_k) (\check{\mathbf{y}}_{k,i} - \boldsymbol{\mu}_{y,k})^T \quad (44c)$$

where

$$\alpha_i = \begin{cases} \frac{\kappa}{L+\kappa} & i = 0 \\ \frac{1}{2} \frac{1}{L+\kappa} & \text{otherwise} \end{cases} \quad (45)$$

- these are plugged into the generalized Gaussian correction-step equations above to complete the correction step


# SPKF commentary

- a big advantage of the SPKF (over the EKF) is that it does not require any analytical derivatives of the motion and observation models
- in fact, the models could even just be black-box software functions
  - no need to express them mathematically!
- there are some additional efficiencies that can be found in the calculations when the observation model becomes:

$$\mathbf{y}_k = \mathbf{g}(\mathbf{x}_k) + \mathbf{n}_k, \quad (46)$$


- there is even an **iterated SPKF** that can be formulated using the idea of **statistical Jacobians**

# SPKF example



- we see that the SPKF does converge to a slightly different place than the IEKF on our simple stereo camera example
- it does not go to the MAP solution...

# SPKF example



- over 1,000,000 we find the ISPKF actually converges quite close to the **mean** of the true posterior, not the **mode**

# Summary

- we investigated a few more alternatives in our taxonomy of filters
- by using Monte Carlo sampling, we came up with the **particle filter**, which can be used for general nonlinear, non-Gaussian problems, but is quite inefficient in most cases
- by using the **sigmapoint transformation**, we came up with the **sigmapoint Kalman filter**, which is only slightly less efficient than the EKF, but is quite a bit more accurate on nonlinear problems
- the choice of EKF, PF, or SPKF depends on how nonlinear and non-Gaussian your problem is